

MAST90083 Assignment 1

Haonan Zhong

Question 1.1

```
library("MASS")
library("ISLR")
suppressMessages(library("glmnet"))
data(Hitters)
# Remove rows with NA in the Salary column
Hitters <- Hitters[!is.na(Hitters$Salary),]
```

Question 1.2

```
# Construct design matrix and response variable
x <- model.matrix(~.-1, data = subset(Hitters, select = -(Salary)))
y <- Hitters$Salary
lambda <- 10^seq(10, -2, length = 100)

# Estimate ridge coefficients for 100 lambda values
coef_estimate <- glmnet(x, y, alpha = 0, lambda = lambda)
# Observe the coefficients for the largest lambda
coef(coef_estimate)[,1]
```

```
##      (Intercept)      AtBat      Hits      HmRun      Runs
## 5.359257e+02 5.443467e-08 1.974589e-07 7.956523e-07 3.339178e-07
##      RBI      Walks      Years      CAtBat      CHits
## 3.527222e-07 4.151323e-07 1.697711e-06 4.673743e-09 1.720071e-08
##      CHmRun      CRuns      CRBI      CWalks      LeagueA
## 1.297171e-07 3.450846e-08 3.561348e-08 3.767877e-08 5.800263e-07
##      LeagueN      DivisionW      PutOuts      Assists      Errors
## -5.800262e-07 -7.807263e-06 2.180288e-08 3.561198e-09 -1.660460e-08
##      NewLeagueN
## -1.152288e-07
```

```
#Observe the coefficients for the smallest lambda
coef(coef_estimate)[,100]
```


```
##      (Intercept)      AtBat      Hits      HmRun      Runs
## 195.91660826 -1.97384734 7.37766330 3.93654704 -2.19869563
##      RBI      Walks      Years      CAtBat      CHits
## -0.91621489 6.20035808 -3.71378518 -0.17510678 0.21135663
##      CHmRun      CRuns      CRBI      CWalks      LeagueA
```

```
##      0.05633262      1.36604022      0.70963872     -0.79581596    -31.80455226
##      LeagueN      DivisionW      PutOuts      Assists      Errors
##      31.60419354    -117.08236876      0.28202517      0.37318700     -3.42404313
##      NewLeagueN
##     -25.99406318
```

As we can see from the output above, coefficients for the largest lambda is much more closer to 0. Which is quite reasonable, as the effect of shrinkage penalty grows as λ increases, and the ridge regression coefficients will get closer to 0.

Question 1.3

```
l2norm <- rep(0, length(lambda))
for (i in 1:100) {
  l2norm[i] <- norm(coef_estimate$beta[,i], type="2")
}
plot(log(lambda), l2norm, xlab = "Logrithm of lambda", ylab = "l2-norm of coefficients")
```




```
# Computing the MSE for each lambda
mse <- rep(0, length(lambda))
for (i in 1:length(lambda)) {
  prediction <- rep(0, length(y))
  for (j in 1:length(y)) {
```

```

    prediction[j] <- t(x[j,]) %*% matrix(coef_estimate$beta[, i])
  }
  mse[i] <- mean((y - prediction)^2)
}

plot(log(lambda), mse, xlab = "Logrithm of lambda", ylab = "Mean Squared Error")

```



We cannot really say anything about the optimal value of λ with l_2 -norm, since it only tells us the size of the coefficients. On the other hand, mean squared error can tell us how accurate the coefficients are in terms of estimating the our response variable, *Salary*.

Question 1.4

```


# Set the seed equal to 10 for random number generator
set.seed(10)
# Sample the training set
n_train <- sample(seq_len(length(y)), size = 131)
x_train <- x[n_train, ]
y_train <- y[n_train]
# Sample the testing set
x_test <- x[-n_train, ]
y_test <- y[-n_train]

# Performing 10-fold cross validation
(train_cv <- cv.glmnet(x_train, y_train, lambda = lambda, type.measure = "mse", alpha = 0))

```

```
##
## Call:  cv.glmnet(x = x_train, y = y_train, lambda = lambda, type.measure = "mse",      alpha = 0)
##
## Measure: Mean-Squared Error
##
##      Lambda Index Measure      SE Nonzero
## min  305.4    63  95669 15586      20
## 1se 2154.4    56 110755 17714      20
```

```
plot(train_cv)
```



As the result above depicted, the model has the lowest mean squared error at 95,669 when $\lambda = 305.4$. Next, we will evaluate the test mean squared error.

```
test_pred <- predict(train_cv$glmnet.fit, train_cv$lambda.min, newx = x_test)
mean((y_test - test_pred)^2)
```

```
## [1] 143265.4
```

The corresponding MSE for $\lambda = 305.4$ on the testing set is 143265.4.

```
# Refit the ridge regression model on the full data set using the lambda chosen by CV
new_model <- glmnet(x, y, alpha = 0, lambda = train_cv$lambda.min)
coef(new_model)
```

```
## 21 x 1 sparse Matrix of class "dgCMatrix"
##              s0
## (Intercept) 26.64638332
## AtBat       0.07272471
## Hits        0.88144658
## HmRun       0.55633590
## Runs        1.07968045
## RBI         0.88278958
## Walks       1.64777011
## Years       1.21365821
## CAtBat      0.01134502
## CHits       0.05845175
## CHmRun      0.41373316
## CRuns       0.11649930
## CRBI        0.12319621
## CWalks      0.04999236
## LeagueA    -15.84232371
## LeagueN     15.84278645
## DivisionW  -80.96578812
## PutOuts     0.16967851
## Assists     0.03076577
## Errors     -1.45580137
## NewLeagueN  4.54951346
```

```
ols_model <- lm(Salary ~ ., data = Hitters)
ols_model$coefficients
```

```
## (Intercept)      AtBat      Hits      HmRun      Runs      RBI
## 163.1035878 -1.9798729  7.5007675  4.3308829 -2.3762100 -1.0449620
##      Walks      Years      CAtBat      CHits      CHmRun      CRuns
##   6.2312863 -3.4890543 -0.1713405  0.1339910 -0.1728611  1.4543049
##      CRBI      CWalks      LeagueN      DivisionW      PutOuts      Assists
##   0.8077088 -0.8115709 62.5994230 -116.8492456  0.2818925  0.3710692
##      Errors      NewLeagueN
##  -3.3607605 -24.7623251
```

As we can see from the output above, the coefficients of the ridge regression model are much smaller compare to the one from ordinary least square model. In some cases, the coefficients in ridge regression are shrinked close to zero, but ridge regression still retains all the variables.


Question 1.5

```
set.seed(10)
(lasso_cv <- cv.glmnet(x_train, y_train, alpha = 1, lambda = lambda, type.measure = "mse"))

##
## Call:  cv.glmnet(x = x_train, y = y_train, lambda = lambda, type.measure = "mse",      alpha = 1)
##
## Measure: Mean-Squared Error
##
##      Lambda Index Measure      SE Nonzero
```

```
## min 18.74    73   95029 14464    7
## 1se 75.65    68  103753 17212    4
```

```
plot(lasso_cv)
```



```
lasso_test_pred <- predict(lasso_cv$glmnet.fit, lasso_cv$lambda.min, newx = x_test)
mean((y_test - lasso_test_pred)^2)
```

```
## [1] 142270.1
```

As the output above suggests, the optimal λ value for lasso regression is 18.74, and the corresponding mean squared error for the testing set is 142270.1.

```
lasso_new_model <- glmnet(x, y, alpha = 1, lambda = lasso_cv$lambda.min)
coef(lasso_new_model)
```

```
## 21 x 1 sparse Matrix of class "dgCMatrix"
##                               s0
## (Intercept) 24.6396990
## AtBat      .
## Hits      1.8499834
## HmRun     .
## Runs      .
## RBI       .
```

```

## Walks      2.1959762
## Years      .
## CAtBat     .
## CHits      .
## CHmRun     .
## CRuns      0.2058514
## CRBI       0.4095910
## CWalks     .
## LeagueA    .
## LeagueN    .
## DivisionW  -99.9733911
## PutOuts    0.2158910
## Assists    .
## Errors     .
## NewLeagueN .

```

Finally, we refitted the lasso regression model using the $\lambda = 18.74$ selected from cross-validation. As we can see, most of the coefficients were shrunk to zero, thus, less important variables are eliminated when penalized, resulted in a sparse model compare to ordinary least square and ridge regression model.